

Deep Learning  
EE-559 - EPFL

---

## Mini-project II

Implementing from scratch a mini deep-learning framework

---

Alexandre CARLIER

[alexandre.carlier@epfl.ch](mailto:alexandre.carlier@epfl.ch)

May, 2018



ÉCOLE POLYTECHNIQUE  
FÉDÉRALE DE LAUSANNE

# 1 Introduction

In this project, we build a small deep learning framework called **Autogradlib** using only PyTorch's **Tensor** operations and the standard math library. The autograd mechanism and the high-level syntax are inspired by PyTorch.

Autogradlib builds a gradient operations directed acyclic graph (DAG) on the fly, which enables to perform efficiently the backprop algorithm and to draw the associated graph using **graphviz** (see figure 1).

## Note on the submission files:

The program can be executed using three different files:

- **test.py**: as it was suggested to submit a single executable file, this self-sufficient file includes all the code. However, as the program is written in an Object-oriented style, it may be easier to read it from the **autogradlib** package.
- **test\_fancy.py**: same as above, except that it also includes methods to draw the gradient operations DAG from the network (with the **graphviz** package).
- **test\_with\_library.py**: same as above except that it uses the **autogradlib** package.

# 2 Implementation details

The code structure is shown in figure 2

## 2.1 The Variable and GradOperator base classes

Like in PyTorch, when building a neural network, we wrap **Tensors** in a **Variable** object. The **Variable** class stores:

- **tensor**: the wrapped **Tensor**
- **grad**: the gradient of the loss w.r.t. the tensor
- **grad\_op**: the **GradOperator** instance that actually performs the backpropagation of gradients

The **GradOperator** class stores:

- **variable**: the wrapped output variable
- **children**: the children nodes in the gradient operations DAG

Thus, when running the following piece of code:

```
x = Variable(FloatTensor(3, 3).normal_())
y = x.t() - x.exp()
z = y.pow(2).sum()
dot = z.draw_graph()
dot.render("model.gv", view=True)
```

it will produce an underlying tree of **GradOperators** that we show in figure 3

When we then call **z.backward()**, the **Variable** object makes a call to **self.grad\_op.pass\_grad(FloatTensor([1]))** which in turn passes recursively the gradient to the **GradOperator** child nodes.

**GradOperator::pass\_grad(gradwrtoutput)** takes as argument the gradient of the loss w.r.t. the output. The method accumulates the current variable's **grad** with this provided gradient and passes to its children the gradient of the loss w.r.t. their respective outputs.




Figure 1: The gradient operations graph associated to the final network. This graph is rendered automatically in Autogradlib with the `draw_graph` method.

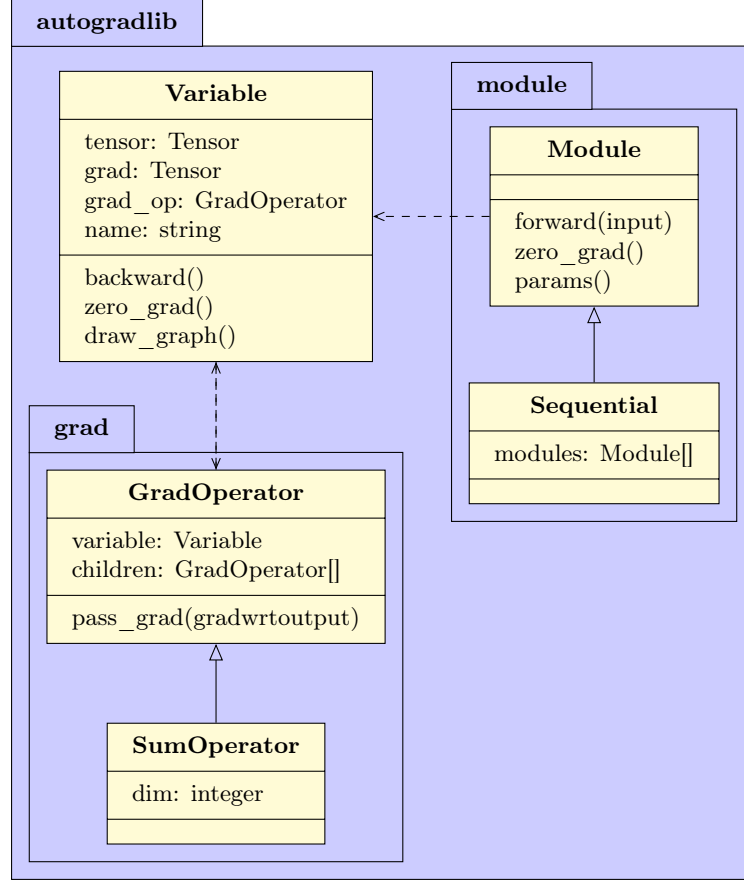


Figure 2: Autogradlib’s code structure including the three main classes **Variable**, **GradOperator** and **Module**, and some classes that inherit from the latter.

## 2.2 The GradOperator subclasses

The **GradOperator** subclasses are responsible of:

- computing the output **Variable** (in the constructor)
- implementing the `pass_grad(self, gradwrtoutput)` method to actually perform the backward pass.

I implemented 11 subclasses (see **grad** package) that enable to perform most of the needed operations in Deep Learning (addition, coefficient-wise multiplication, matrix multiplication, transposition, sum along a specified axis...). To define `pass_grad`, we use the chain rule. See for instance `grad/AddOperator.py` which takes as input two **Variables**  $a$  and  $b$ , stores, the output `Variable(a.tensor + b.tensor)` and passes the unchanged `gradwrtoutput` to the two children in the backward pass. Indeed, let  $f = a + b$ . Then

$$\frac{\partial l}{\partial a} = \frac{\partial f}{\partial a} \frac{\partial l}{\partial f} = \frac{\partial l}{\partial f} \text{ and } \frac{\partial l}{\partial b} = \frac{\partial f}{\partial b} \frac{\partial l}{\partial f} = \frac{\partial l}{\partial f}$$

As another example, see `grad/ExpOperator.py`. If we define  $B = \exp(A)$ , where  $A$  and  $B$  are tensors of same size and `exp` is applied coefficient-wise, then we have:

$$\frac{\partial l}{\partial A_{ijk}} = \frac{\partial B_{ijk}}{\partial A_{ijk}} \times \frac{\partial l}{\partial B_{ijk}} = \exp A_{ijk} \times \frac{\partial l}{\partial B_{ijk}}$$

and thus:

$$\frac{\partial l}{\partial A} = \exp A \odot \frac{\partial l}{\partial B}$$

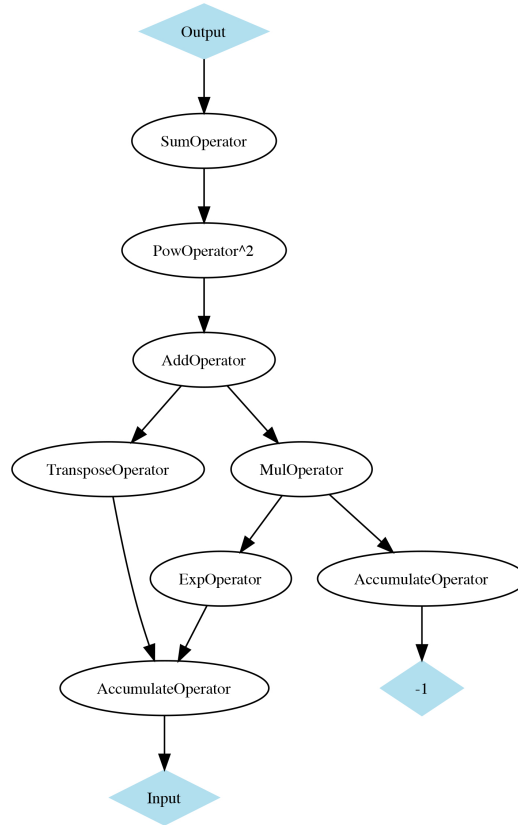


Figure 3: A simple DAG representing the computation `(x.t() - x.exp()).pow(2).sum()`

### 2.3 The Module class

This class is the easiest one to implement. It defines:

- a virtual method `forward(input)` that all its subclasses must implement
- the method `zero_grad()` that re-initializes the gradients of the `Module`'s parameters
- `params()`: automatically computes the list of the `Module`'s parameters by checking its attributes. The supported parameter types are `Variable`, `Module` and list of `Variable` or `Module`

Once the `Module` class is implemented, defining subclasses like `ReLU` or `Sigmoid` simply consists of overriding the `forward(self, x)` method, generally in one line of code! (parameters are computed automatically by `Module` and the backward pass is handled by the corresponding `Variable`).

The more tricky part was `Softmax` since it required to implement two new `Variable` methods: `sum` along a specified dimension and `repeat` in order to multiply two similar sized tensors (since I didn't implement broadcasting like in PyTorch).

Finally, in order to deal with the vanishing gradients problem and control the variance of both activations and gradients of the loss w.r.t. the activations, we have to make sure to initialize the weights using e.g. Xavier initialization (see implementation in the `Linear` class).

## 3 Results

### 3.1 Evaluation

We train a network with two input units, two output units and three hidden layers of 25 units.

```
Sequential(
  0: Linear(in_dim=2, out_dim=25)
```

```

1: ReLU
2: Linear(in_dim=25, out_dim=25)
3: ReLU
4: Linear(in_dim=25, out_dim=25)
5: ReLU
6: Linear(in_dim=25, out_dim=2)
7: Softmax(dim=1)
)

```

Training this model during 500 epochs (repeated 20 times to also compute the variance) gives the following results:

	Loss (median $\pm$ std.)	Accuracy (median $\pm$ std.)
Train dataset	0.0228 $\pm$ 0.0023	0.986 $\pm$ 0.0034
Test dataset	0.0254 $\pm$ 0.0028	0.983 $\pm$ 0.0055

The learning curves are showed in figure 4.




Figure 4: Learning curves of the network

### 3.2 Performance comparison with PyTorch

We run the training of the same model 20 times on a MacBook Pro (2.7 Ghz Intel Core i5) and compute the mean and median execution durations using Autogradlib and PyTorch.

Library	Mean time	Median time
PyTorch	0.804s	0.779s
Autogradlib	1.961s	1.920s

We observe that Autogradlib is only a bit more than twice as slow as PyTorch. However, we must notice that the heavy tensor operations are actually handled by PyTorch's **Tensor** in both cases. Thus, the difference in execution time is mainly due to the fact that Autogradlib's gradient operations graph is computed in Python, while PyTorch implements it in C from what I understood.

## 4 Conclusion

In this project, we show that by simply using the chain rule and a DAG structure, we can build a relatively powerful small deep learning framework.